



Sebastian Huynh
Project 2: SQL Database
October 31, 2023
CIS 3050-04
Fall 2023

Table of Contents

1. Table of Contents.....	1
2. Statement of Academic Honesty.....	2
3. Introduction.....	3
4. Project Descriptions.....	4
5. Test Queries.....	5-27
6. Results & Discussions.....	28
7. Lessons Learned.....	28
8. Conclusion.....	29
9. Cited References.....	30

Statement of Academic Honesty

My name is: Sebastian Huynh, I declare that, except where fully referenced, no aspect of this project has been copied from any other source. I understand that any act of Academic Dishonesty such as plagiarism or collusion may result in serious offense and punishments. I promise not to lie about my academic work, to cheat, or to steal the words or ideas of others, nor will I help fellow students to violate the Code of Academic Honesty.

Name: Sebastian Huynh

Date: Oct 31, 2023

Introduction

Previously in project 1, we explored the integrity and structure of relational databases using visual data modeling. By highlighting meaningful connections between entities, companies are able to organize large amounts of input for advanced purposes such as enhanced decision-making and predictive analysis (that go beyond just record keeping). Proceeding to project 2, the focus shifts from database representations and toward the creation of an actual, working database that adheres to desired data model guidelines. Structured query language has become one of the most popular methods for communication taking place to and from databases and data-related applications (Hoffer, 2018). It is one of the standardized ways of creating and manipulating a working relational database. Furthermore, there are three subcategories for writing SQL which range from data definition, data manipulation, and data control, with the first two being essential for development in a testing-only environment. Data definition SQL statements “specify a database's structure” (Panigrahi, 2023) while data manipulation SQL statements alter or retrieve data from an already established database. Project 2 consists primarily of such DDL and DML solutions.

Project Description

Software Required: Microsoft SQL Server Management Studio, SQL Server Express or Developer, and Windows OS

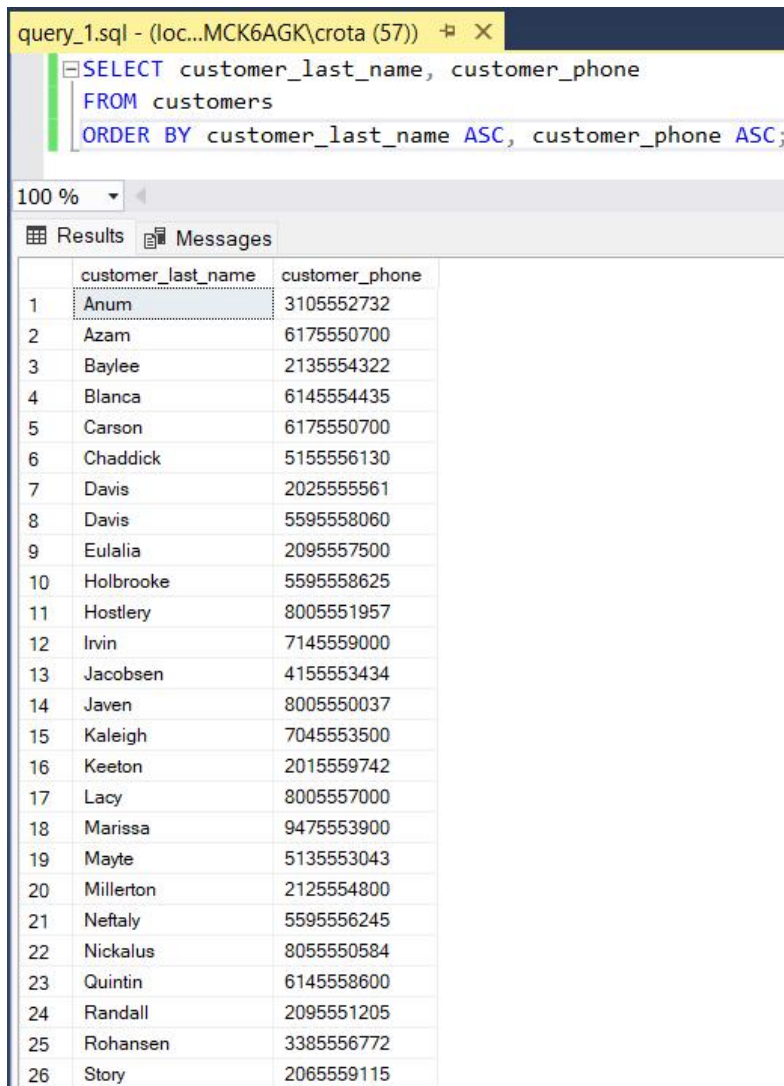
- Management Studio is the encompassing program and integrated environment that provides user-friendly tools as a bridge to databases hosted by SQL Server
- SQL Server is Microsoft's relational database system software that holds all database information and can be accessed by other programs like management studio.
- Management Studio and SQL Server can only be installed on Windows, but Linux users may also use SQL Server if they configure it.

Project Objective: Demonstrate the ability to develop a relational database using MSSQL for the order fulfillment process of a company.

- With input values given, create six tables, one for each entity using DDL statements. Establish relationships between entities using constraints.
- Run DML statements that will insert the input values into the empty tables then use select statements to bring back filtered results according to project instructions. Also, demonstrate the ability to insert and delete data from specified tables.

Query 1: Write a query that displays a list of all customers showing the customer's last name, and phone number. Sort the results by customer last name, then phone number.

```
SELECT customer_last_name, customer_phone FROM customers ORDER BY  
customer_last_name ASC, customer_phone ASC;
```



The screenshot shows a SQL query editor window titled 'query_1.sql - (loc...MCK6AGK\crota (57))'. The query is: `SELECT customer_last_name, customer_phone FROM customers ORDER BY customer_last_name ASC, customer_phone ASC;`. Below the query editor, the 'Results' tab is selected, displaying a table with 26 rows. The table has two columns: 'customer_last_name' and 'customer_phone'. The rows are sorted by last name (A-Z) and then by phone number (smallest to largest).

	customer_last_name	customer_phone
1	Anum	3105552732
2	Azam	6175550700
3	Baylee	2135554322
4	Blanca	6145554435
5	Carson	6175550700
6	Chaddick	5155556130
7	Davis	2025555561
8	Davis	5595558060
9	Eulalia	2095557500
10	Holbrooke	5595558625
11	Hostlery	8005551957
12	Irvin	7145559000
13	Jacobsen	4155553434
14	Javen	8005550037
15	Kaleigh	7045553500
16	Keeton	2015559742
17	Lacy	8005557000
18	Marissa	9475553900
19	Mayte	5135553043
20	Millerton	2125554800
21	Neftaly	5595556245
22	Nickalus	8055550584
23	Quintin	6145558600
24	Randall	2095551205
25	Rohansen	3385556772
26	Story	2065559115

Explanation: Showing only two column fields, customer's last name and their corresponding phone number from the customer's table. Resulting rows are displayed from A-Z according to their last names and then smallest to largest according to their phone numbers. This is sorted by default ascending order.

Query 2: Write a query that displays a list of all customers showing the customer's first name, last name, phone number and fax. Sort the results by customer fax number in ascending order.

```
SELECT customer_first_name, customer_last_name, customer_phone,
customer_fax FROM customers ORDER BY customer_fax ASC;
```

query_2.sql - (loc...MCK6AGK\crota (53)) query_1.sql - (loc...MCK6AGK\crota (52))

```
SELECT customer_first_name, customer_last_name, customer_phone, customer_fax
FROM customers
ORDER BY customer_fax ASC;
```

100 %

Results Messages

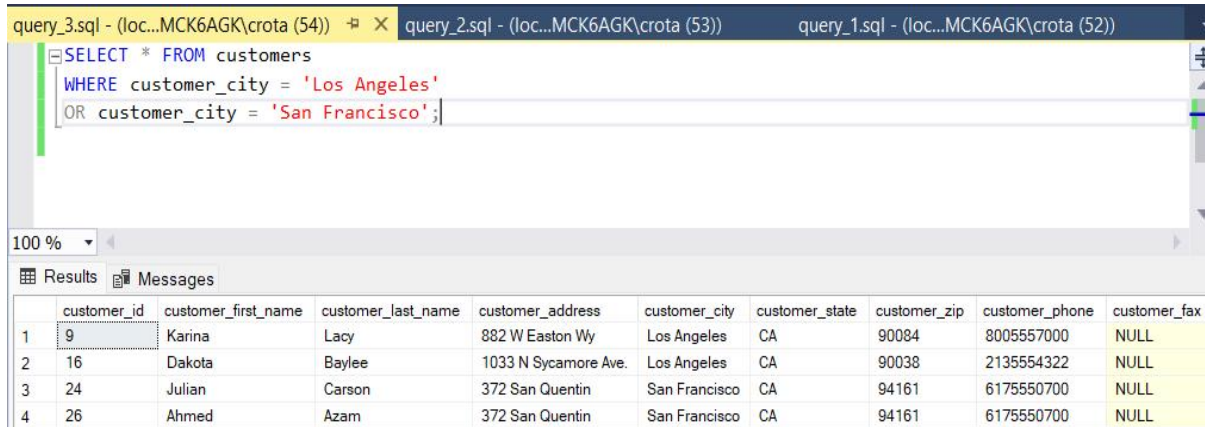
	customer_first_name	customer_last_name	customer_phone	customer_fax
1	Johnathon	Millerton	2125554800	NULL
2	Mikayla	Davis	2025555561	NULL
3	Kendall	Mayte	5135553043	NULL
4	Derek	Chaddick	5155556130	NULL
5	Deborah	Davis	5595558060	NULL
6	Karina	Lacy	8005557000	NULL
7	Anders	Rohansen	3385556772	NULL
8	Thalia	Neftaly	5595556245	NULL
9	Gonzalo	Keeton	2015559742	NULL
10	Ania	Irvin	7145559000	NULL
11	Dakota	Baylee	2135554322	NULL
12	Samuel	Jacobsen	4155553434	NULL
13	Justin	Javen	8005550037	NULL
14	Kyle	Marissa	9475553900	NULL
15	Erick	Kaleigh	7045553500	NULL
16	Trisha	Anum	3105552732	NULL
17	Julian	Carson	6175550700	NULL
18	Kirsten	Story	2065559115	NULL
19	Ahmed	Azam	6175550700	NULL
20	Kurt	Nickalus	8055550584	055556689
21	Kelsey	Eulalia	2095557500	2095551302
22	Yash	Randall	2095551205	2095552262
23	Rashad	Holbrooke	5595558625	5595558495
24	Korah	Blanca	6145554435	6145553928
25	Marvin	Quintin	6145558600	6145557580
26	Kaitlin	Hostlery	8005551957	8005552826

Query executed successfully.

Explanation: Showing only four column fields, a customer's first and last name and then their corresponding phone number and fax number. These fields are derived from the customers table and the resulting rows are sorted from least to greatest (ascending) fax numbers with the least starting from no value or null.

Query 3: Write a query that displays all the customers from San Francisco or Los Angeles in the “Customers” table.

```
SELECT * FROM customers WHERE customer_city = 'Los Angeles' OR  
customer_city = 'San Francisco';
```



The screenshot shows a SQL query editor with three tabs: query_3.sql, query_2.sql, and query_1.sql. The active tab, query_3.sql, contains the following SQL query:

```
SELECT * FROM customers  
WHERE customer_city = 'Los Angeles'  
OR customer_city = 'San Francisco';
```

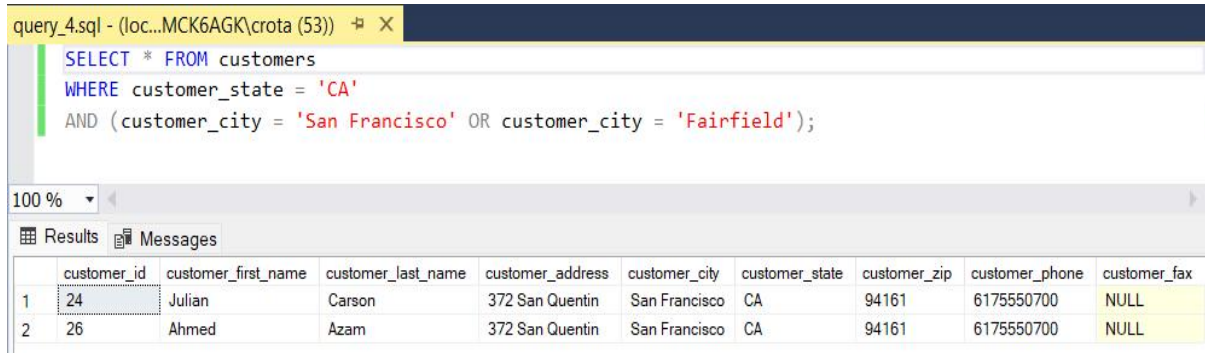
Below the query editor, the 'Results' tab is selected, displaying a table with 10 columns and 4 rows of data. The columns are: customer_id, customer_first_name, customer_last_name, customer_address, customer_city, customer_state, customer_zip, customer_phone, and customer_fax. The rows represent customers from Los Angeles and San Francisco.

	customer_id	customer_first_name	customer_last_name	customer_address	customer_city	customer_state	customer_zip	customer_phone	customer_fax
1	9	Karina	Lacy	882 W Easton Wy	Los Angeles	CA	90084	8005557000	NULL
2	16	Dakota	Baylee	1033 N Sycamore Ave.	Los Angeles	CA	90038	2135554322	NULL
3	24	Julian	Carson	372 San Quentin	San Francisco	CA	94161	6175550700	NULL
4	26	Ahmed	Azam	372 San Quentin	San Francisco	CA	94161	6175550700	NULL

Explanation: This statement shows all available column fields from the customers table for each row having the condition that the customer is either from the city of Los Angeles or San Francisco being met. A customer cannot be part from more than one city according to the field specifications so a customer instance will be returned if just one of these two conditions are met (Los Angeles being the first and San Francisco being the second).

Query 4: Write a query that displays all the customers from the state of California and live in San Francisco or Fairfield.

```
SELECT * FROM customers WHERE customer_state = 'CA'
AND (customer_city = 'San Francisco' OR customer_city = 'Fairfield');
```



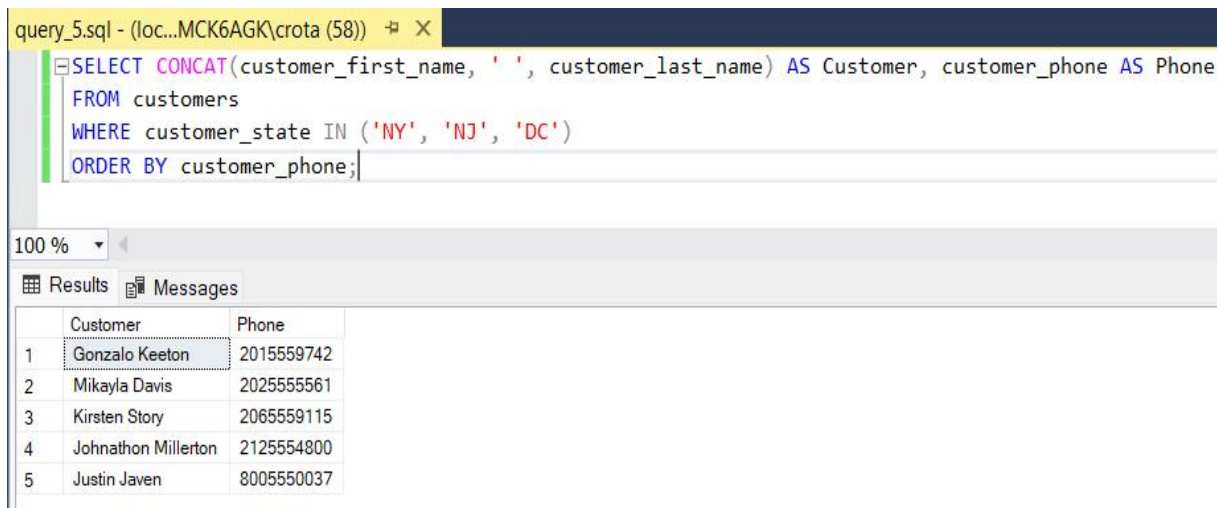
The screenshot shows a SQL query editor window titled 'query_4.sql'. The query is: `SELECT * FROM customers WHERE customer_state = 'CA' AND (customer_city = 'San Francisco' OR customer_city = 'Fairfield');`. Below the query editor, there is a 'Results' tab showing a table with 10 columns: customer_id, customer_first_name, customer_last_name, customer_address, customer_city, customer_state, customer_zip, customer_phone, and customer_fax. The table contains two rows of data.

	customer_id	customer_first_name	customer_last_name	customer_address	customer_city	customer_state	customer_zip	customer_phone	customer_fax
1	24	Julian	Carson	372 San Quentin	San Francisco	CA	94161	6175550700	NULL
2	26	Ahmed	Azam	372 San Quentin	San Francisco	CA	94161	6175550700	NULL

Explanation: This statement shows all available column fields from the customers table, returning every row which meets the condition that the customer is living in either the city of San Francisco or the city of Fairfield, located only in the state of California. This means that any customers who live in a different Fairfield for example outside of California is not returned in the results.

Query 5: Write a query that displays each customer name as a single field in the format “firstname lastname” with a heading of Customer, along with their phone number with a heading of Phone. Use the IN operator to only display customers in New York, New Jersey, or Washington D.C. Sort the results by phone number.

```
SELECT CONCAT(customer_first_name, ' ', customer_last_name) AS Customer, customer_phone AS Phone FROM customers WHERE customer_state IN ('NY', 'NJ', 'DC') ORDER BY customer_phone;
```



The screenshot shows a SQL query editor with the following query:

```
SELECT CONCAT(customer_first_name, ' ', customer_last_name) AS Customer, customer_phone AS Phone
FROM customers
WHERE customer_state IN ('NY', 'NJ', 'DC')
ORDER BY customer_phone;
```

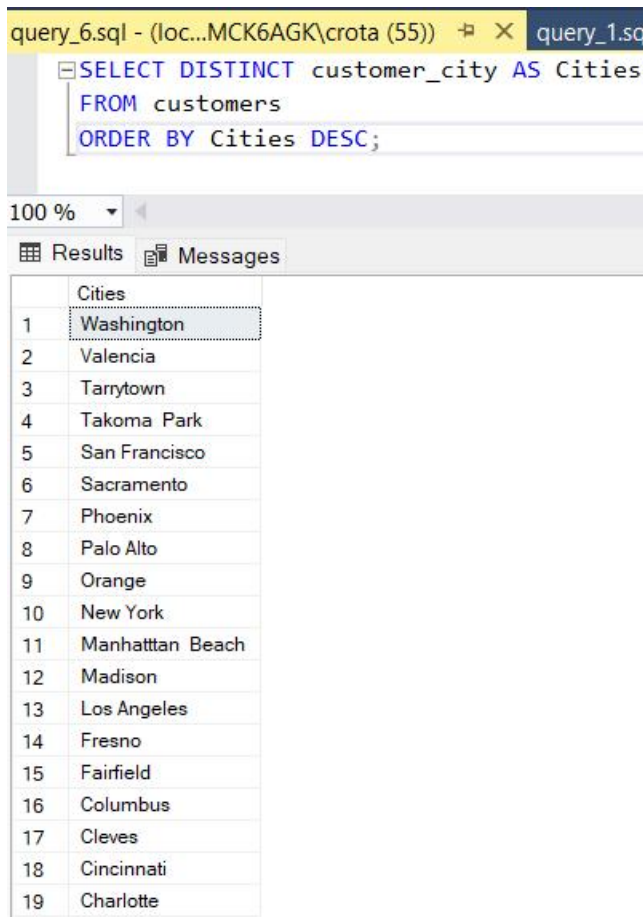
Below the query editor, the results are displayed in a table with two columns: Customer and Phone. The results are sorted by phone number in ascending order.

	Customer	Phone
1	Gonzalo Keeton	2015559742
2	Mikayla Davis	2025555561
3	Kirsten Story	2065559115
4	Johnathon Millerton	2125554800
5	Justin Javen	8005550037

Explanation: The first field returned is a string combination of customers’ first names followed by a space and then their last name for each returned row. The alias name for this first field is ‘Customer’. For the second field, the column containing customer phone numbers is returned under the alias ‘Phone’. Everything is derived from the customers table. Everything is derived from the customers table and only customers who live in either New York, New Jersey, or Washington DC are shown in ascending (least to greatest) order of their corresponding phone numbers.

Query 6: Write a query that will list all the cities that have customers with a heading of Cities. Only list each city once (no duplicates) and sort in descending alphabetical order.

```
SELECT DISTINCT customer_city AS Cities FROM customers ORDER BY Cities DESC;
```



The screenshot shows a SQL query editor window with the following query:

```
SELECT DISTINCT customer_city AS Cities
FROM customers
ORDER BY Cities DESC;
```

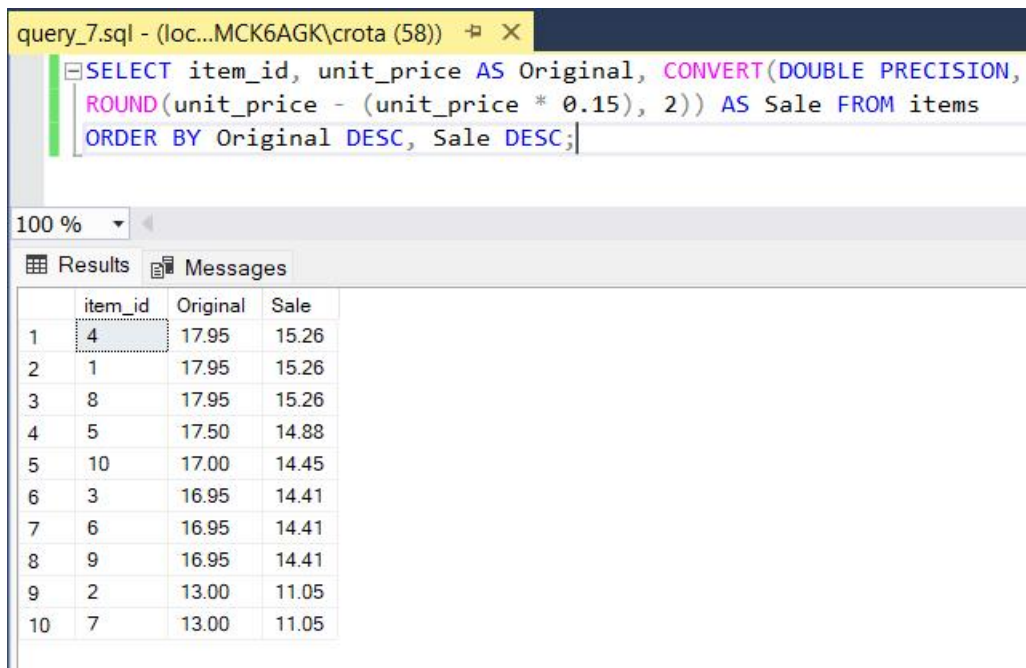
Below the query editor, the 'Results' tab is selected, displaying a table with 19 rows. The table has a single column named 'Cities'. The results are sorted in descending alphabetical order.

	Cities
1	Washington
2	Valencia
3	Tarrytown
4	Takoma Park
5	San Francisco
6	Sacramento
7	Phoenix
8	Palo Alto
9	Orange
10	New York
11	Manhattan Beach
12	Madison
13	Los Angeles
14	Fresno
15	Fairfield
16	Columbus
17	Cleves
18	Cincinnati
19	Charlotte

Explanation: This statement returns only one field with the alias name 'Cities' because we only want to see the list of possible cities a customer in our database could be from. By using the distinct descriptor any city that appears will only show up once in the results even if they appear in the tables multiple times. It is derived from the customers table.

Query 7: Write a query that displays the title of each item along with the price (with a heading of Original) and a calculated field reflecting the price with a 15% discount (with a heading of Sale). Display the sale price with two decimal places using the ROUND function. Sort by price from highest to lowest.

```
SELECT item_id, unit_price AS Original, CONVERT(DOUBLE PRECISION,  
ROUND(unit_price - (unit_price * 0.15), 2)) AS Sale FROM items ORDER BY  
Original DESC, Sale DESC;
```



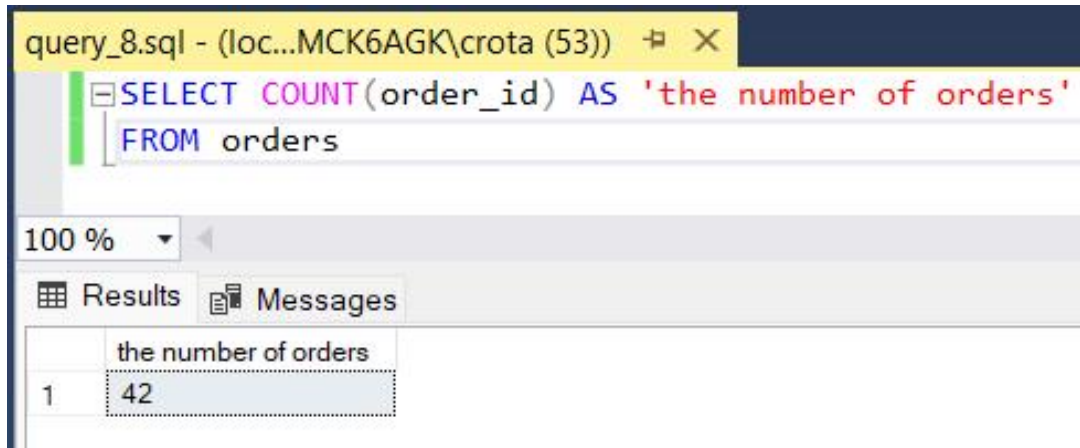
The screenshot shows a SQL query editor window titled 'query_7.sql'. The query is: `SELECT item_id, unit_price AS Original, CONVERT(DOUBLE PRECISION, ROUND(unit_price - (unit_price * 0.15), 2)) AS Sale FROM items ORDER BY Original DESC, Sale DESC;`. Below the query editor, the 'Results' tab is active, displaying a table with 10 rows. The table has four columns: 'item_id', 'Original', and 'Sale'. The 'item_id' column contains values 4, 1, 8, 5, 10, 3, 6, 9, 2, 7. The 'Original' column contains values 17.95, 17.95, 17.95, 17.50, 17.00, 16.95, 16.95, 16.95, 13.00, 13.00. The 'Sale' column contains values 15.26, 15.26, 15.26, 14.88, 14.45, 14.41, 14.41, 14.41, 11.05, 11.05. The first row is highlighted.

	item_id	Original	Sale
1	4	17.95	15.26
2	1	17.95	15.26
3	8	17.95	15.26
4	5	17.50	14.88
5	10	17.00	14.45
6	3	16.95	14.41
7	6	16.95	14.41
8	9	16.95	14.41
9	2	13.00	11.05
10	7	13.00	11.05

Explanation: In the results, the first two columns are originally from the table of items and are used to calculate the third column which is the new price after a 15% discount had been applied to each item. To get the discounted price, we take the original price and subtract it from 15% of the original price. This filtered query allows us to see how much of a difference the original and sales prices are. The round function allows the decimal number to format to cents and double precision is specified to truncate any leading zeros after the 2nd decimal place.

Query 8: Write a query that displays the number of orders.

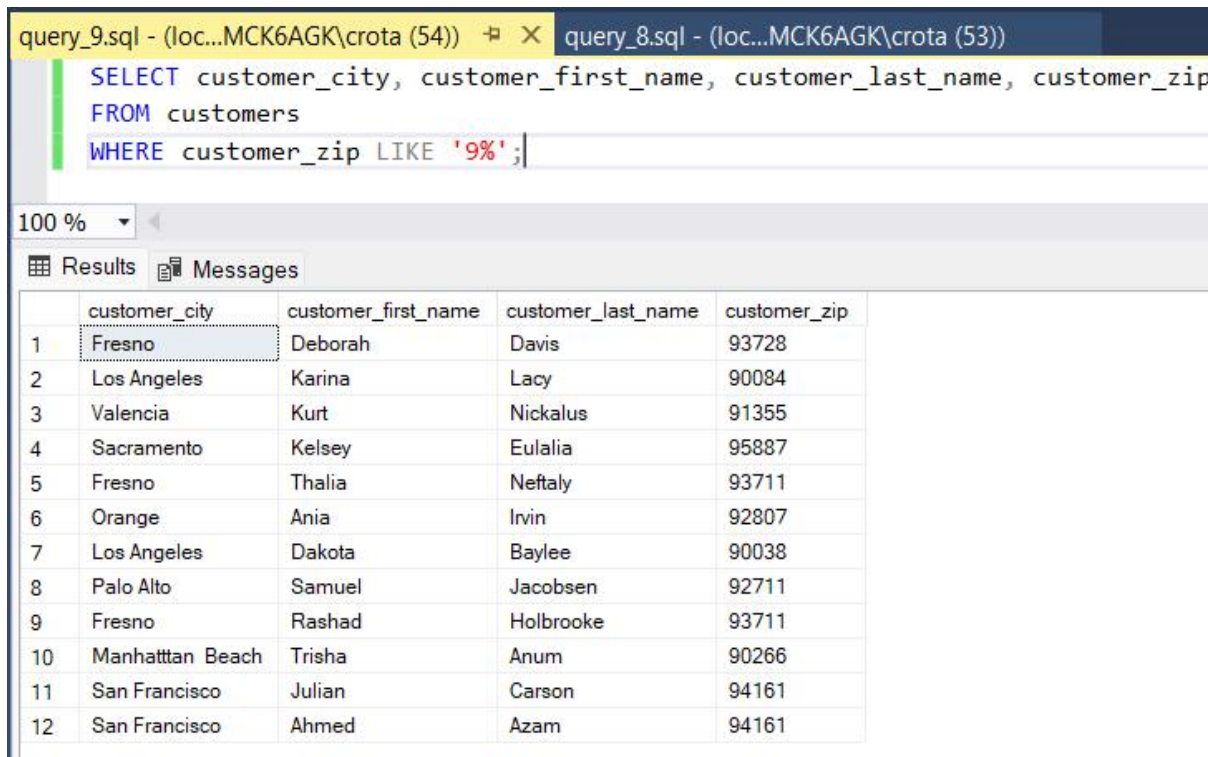
```
SELECT COUNT(order_id) AS 'the number of orders'  
FROM orders
```



Explanation: From the orders table, we are interested the column field of order_id to select unique orders. This is important as we don't want to count the quantity of each order but only how many different orders there are in total. Using the count function we can aggregate this amount.

Query 9: Write a query that displays the customer city, first name, last name, and zip code from the customer's table. Use the LIKE operator to only display customers that reside in any zip code beginning with 9.

```
SELECT customer_city, customer_first_name, customer_last_name, customer_zip  
FROM customers WHERE customer_zip LIKE '9%';
```



The screenshot shows a SQL query editor with two tabs: 'query_9.sql' and 'query_8.sql'. The 'query_9.sql' tab is active, displaying the following SQL query:

```
SELECT customer_city, customer_first_name, customer_last_name, customer_zip  
FROM customers  
WHERE customer_zip LIKE '9%';
```

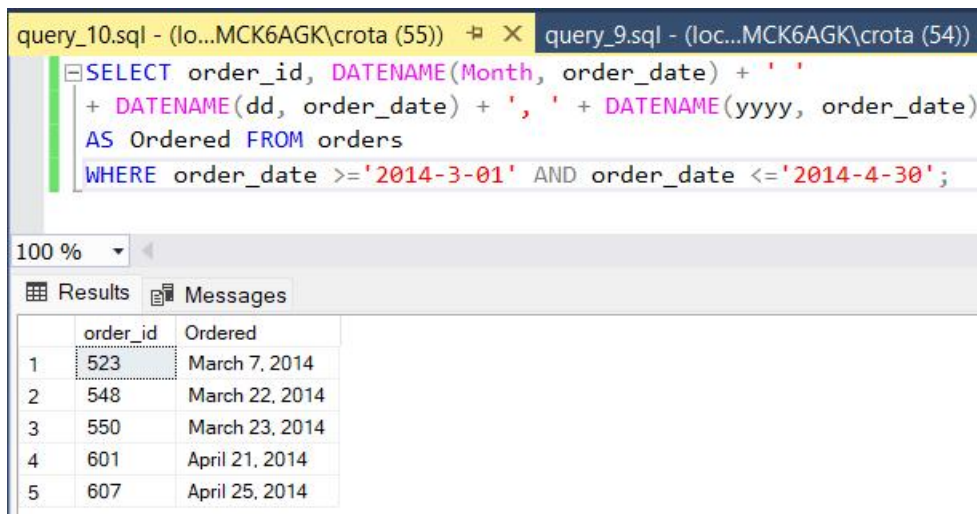
Below the query editor, the 'Results' tab is selected, showing a table with 12 rows and 5 columns. The columns are 'customer_city', 'customer_first_name', 'customer_last_name', and 'customer_zip'. The rows represent customers whose zip codes begin with 9.

	customer_city	customer_first_name	customer_last_name	customer_zip
1	Fresno	Deborah	Davis	93728
2	Los Angeles	Karina	Lacy	90084
3	Valencia	Kurt	Nickalus	91355
4	Sacramento	Kelsey	Eulalia	95887
5	Fresno	Thalia	Neftaly	93711
6	Orange	Ania	Irvin	92807
7	Los Angeles	Dakota	Baylee	90038
8	Palo Alto	Samuel	Jacobsen	92711
9	Fresno	Rashad	Holbrooke	93711
10	Manhattan Beach	Trisha	Anum	90266
11	San Francisco	Julian	Carson	94161
12	San Francisco	Ahmed	Azam	94161

Explanation: Selecting four customer-related column fields from the customers table. In this query, the only records/rows that return are for customers who have zip codes beginning with the number 9. To accommodate for the different numbers after the initial 9 we use the percent sign when putting the condition for customer_zip.

Query 10: Write a query that displays the order id and order date for any orders placed from March 1, 2014 through April 30, 2014. Do this WITHOUT using the BETWEEN clauses. Format the date field as Month dd, yyyy and use a heading of “Ordered”.

```
SELECT order_id, DATENAME(Month, order_date) + ' ' + DATENAME(dd,
order_date) + ', ' + DATENAME(yyyy, order_date) AS Ordered FROM orders
WHERE order_date >='2014-3-01' AND order_date <='2014-4-30';
```



The screenshot shows a SQL query editor with two tabs: 'query_10.sql' and 'query_9.sql'. The 'query_10.sql' tab is active, displaying the following SQL query:

```
SELECT order_id, DATENAME(Month, order_date) + ' ' + DATENAME(dd,
+ DATENAME(yyyy, order_date) AS Ordered FROM orders
WHERE order_date >='2014-3-01' AND order_date <='2014-4-30';
```

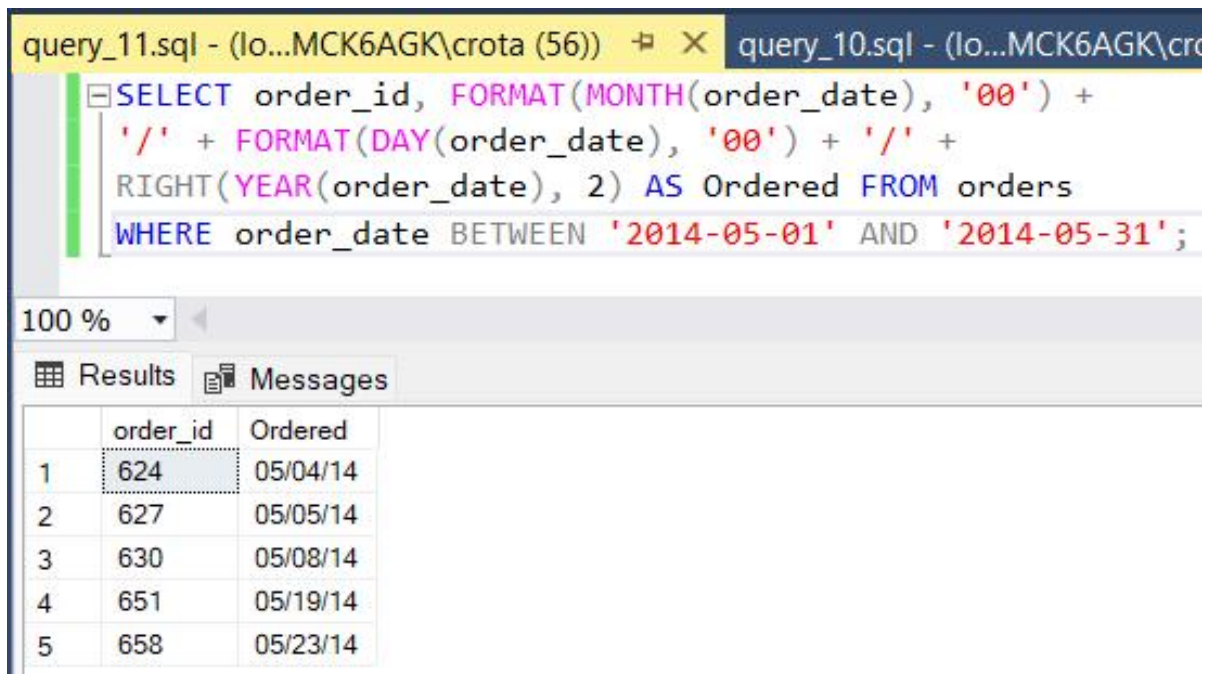
Below the query editor, there is a 'Results' tab showing the output of the query. The results are displayed in a table with two columns: 'order_id' and 'Ordered'.

	order_id	Ordered
1	523	March 7, 2014
2	548	March 22, 2014
3	550	March 23, 2014
4	601	April 21, 2014
5	607	April 25, 2014

Explanation: By not using the between clauses, the alternative is by using comparison operators to create a range from start date and end date. For this query we select the order id field column from orders and then concatenate string for the second returned field using the addition operator. The date name function returns the specific name or number of a date given the time period (month, day, year, etc).

Query 11: Write a query that displays the order id and order date for any orders placed during the month of May 2014. Do this using the BETWEEN clauses. Format the date field as mm/dd/yy and use a heading of “Ordered”.

```
SELECT order_id, FORMAT(MONTH(order_date), '00') + '/' +  
FORMAT(DAY(order_date), '00') + '/' + RIGHT(YEAR(order_date), 2) AS  
Ordered FROM orders WHERE order_date BETWEEN '2014-05-01' AND '2014-  
05-31';
```



The screenshot shows a SQL query editor with two tabs: 'query_11.sql' and 'query_10.sql'. The 'query_11.sql' tab is active, displaying the following SQL query:

```
SELECT order_id, FORMAT(MONTH(order_date), '00') +  
'/' + FORMAT(DAY(order_date), '00') + '/' +  
RIGHT(YEAR(order_date), 2) AS Ordered FROM orders  
WHERE order_date BETWEEN '2014-05-01' AND '2014-05-31';
```

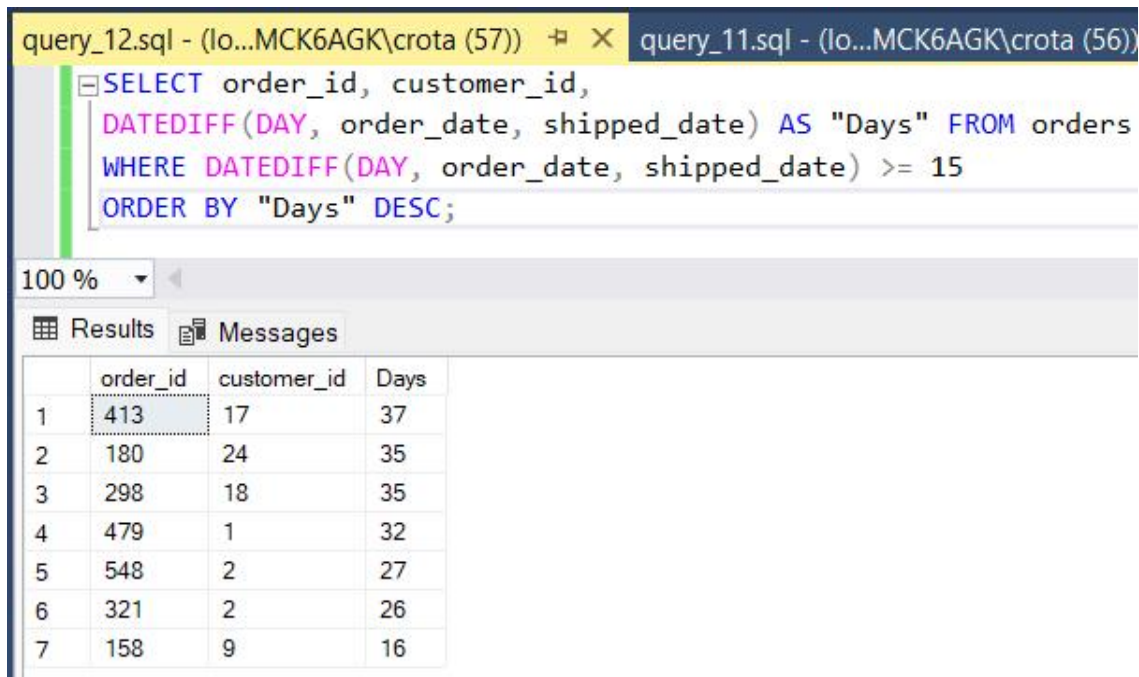
Below the query editor, there is a 'Results' tab showing the output of the query. The results are displayed in a table with two columns: 'order_id' and 'Ordered'. The table contains five rows of data:

	order_id	Ordered
1	624	05/04/14
2	627	05/05/14
3	630	05/08/14
4	651	05/19/14
5	658	05/23/14

Explanation: Two columns are returned using data from the orders table, with the first being the identifying column for each order and the second being a string formatted date column of the order date corresponding to each order. The format function returns 2-digit places for the month and day, and using the right clause allows us to show only the last two digits of the year. Using the between clause constrains dates from the first day of May to the last day of May from 2014 and is the condition for an order to be returned.

Query 12: Write a query which displays the order id, customer id, and the number of days between the order date and the ship date (use the DATEDIFF function). Name this column “Days” and sort by highest to lowest number of days. Only display orders where this result is 15 days or more.

```
SELECT order_id, customer_id,  
DATEDIFF(DAY, order_date, shipped_date) AS "Days" FROM orders  
WHERE DATEDIFF(DAY, order_date, shipped_date) >= 15  
ORDER BY "Days" DESC;
```



The screenshot shows a SQL query editor with two tabs: 'query_12.sql' and 'query_11.sql'. The 'query_12.sql' tab is active, displaying the following SQL query:

```
SELECT order_id, customer_id,  
DATEDIFF(DAY, order_date, shipped_date) AS "Days" FROM orders  
WHERE DATEDIFF(DAY, order_date, shipped_date) >= 15  
ORDER BY "Days" DESC;
```

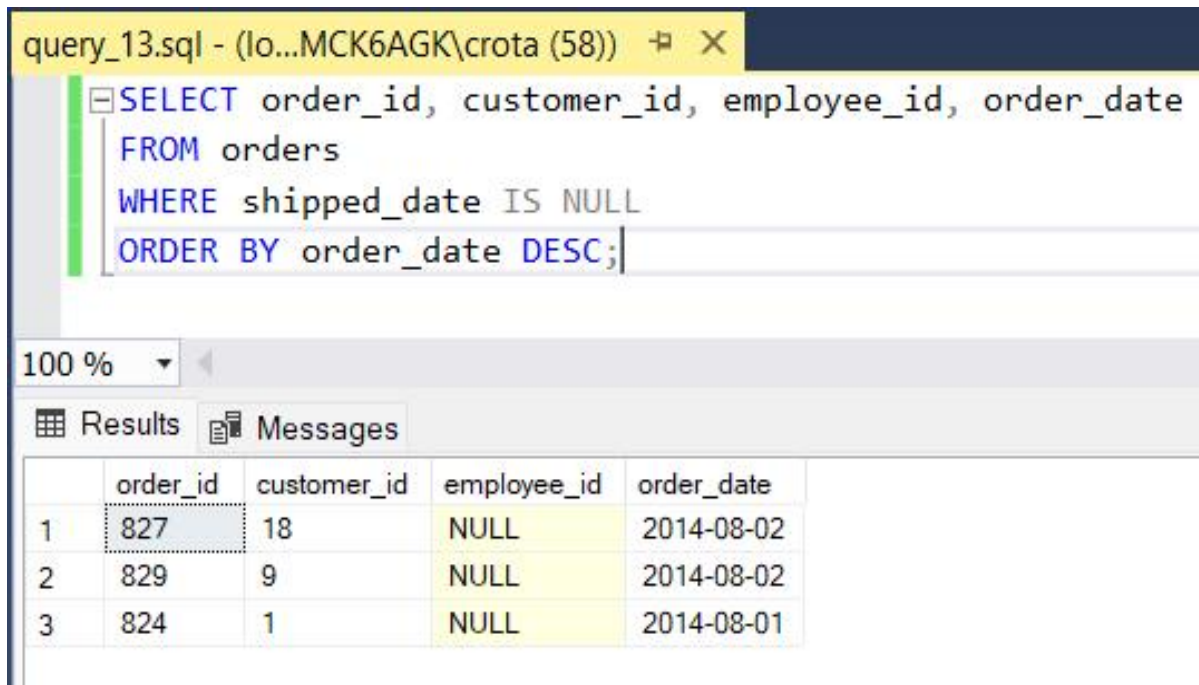
Below the query editor, there is a 'Results' tab showing the output of the query. The results are displayed in a table with the following columns: 'order_id', 'customer_id', and 'Days'. The table contains 7 rows of data, sorted by 'Days' in descending order.

	order_id	customer_id	Days
1	413	17	37
2	180	24	35
3	298	18	35
4	479	1	32
5	548	2	27
6	321	2	26
7	158	9	16

Explanation: From the order tables we are interested in seeing the order and customer combination identified by the id fields and then a calculated field to see the lead time aka the time it took from the day an order was placed to the day it was then shipped. Using the datediff function allows us to show the number difference between two given dates after specifying the time unit (day, month, year, etc). The resulting rows are then sorted from most days to least days it took.

Query 13: Write a query which displays the order id, customer id, employee id, and order date for all orders that have NOT been shipped, sorted by order date with the most recent order at the top.

```
SELECT order_id, customer_id, employee_id, order_date
FROM orders WHERE shipped_date IS NULL ORDER BY order_date DESC;
```



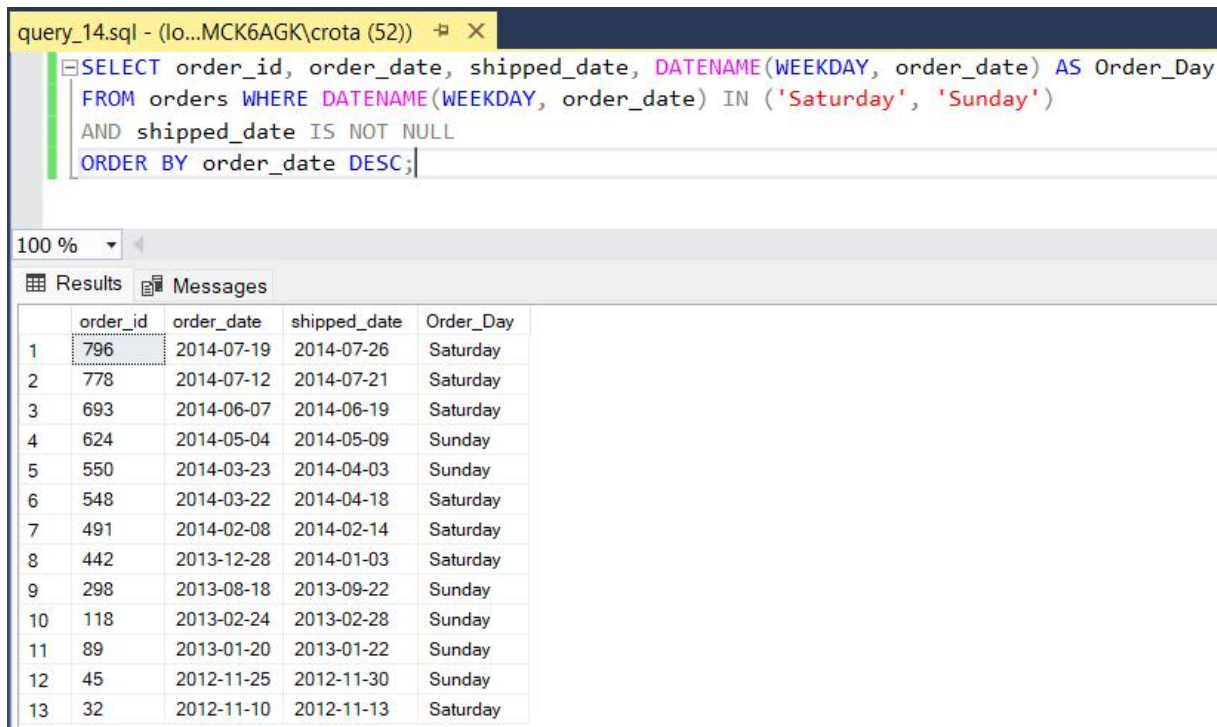
The screenshot shows a SQL query editor window titled 'query_13.sql - (lo...MCK6AGK\crota (58))'. The query is: `SELECT order_id, customer_id, employee_id, order_date FROM orders WHERE shipped_date IS NULL ORDER BY order_date DESC;`. Below the query editor, there is a 'Results' tab showing a table with 5 columns: 'order_id', 'customer_id', 'employee_id', and 'order_date'. The table contains 3 rows of data, sorted by 'order_date' in descending order. The first row has order_id 827, customer_id 18, employee_id NULL, and order_date 2014-08-02. The second row has order_id 829, customer_id 9, employee_id NULL, and order_date 2014-08-02. The third row has order_id 824, customer_id 1, employee_id NULL, and order_date 2014-08-01.

	order_id	customer_id	employee_id	order_date
1	827	18	NULL	2014-08-02
2	829	9	NULL	2014-08-02
3	824	1	NULL	2014-08-01

Explanation: This query is to return any orders that still need to be shipped. The column fields we are interested in is the order id to know which order specifically we need to know about, and also the customer the order belongs to and the employee which may be responsible for the order as well as the date the order was initially placed. All of these columns have been derived from the orders table. The condition for any orders/rows that are returned is when the value in shipping date does not exist aka null. Then we order the rows from most recent to later order date.

Query 14: The Marketing Department has requested a new report of shipped orders for which the order was placed on either a Saturday or a Sunday. Write a query which displays the order id, order date, shipped date, along with a calculated column labeled “Order_Day” showing the day of the week the order was placed (use the DAYNAME function). Only display orders that have shipped and were placed on a Saturday or Sunday. Sort by order date with most recent orders at the top.

```
SELECT order_id, order_date, shipped_date, DATENAME(WEEKDAY,
order_date) AS Order_Day FROM orders WHERE DATENAME(WEEKDAY,
order_date) IN ('Saturday', 'Sunday') AND shipped_date IS NOT NULL ORDER
BY order_date DESC;
```



The screenshot shows a SQL query editor with the following query:

```
SELECT order_id, order_date, shipped_date, DATENAME(WEEKDAY, order_date) AS Order_Day
FROM orders WHERE DATENAME(WEEKDAY, order_date) IN ('Saturday', 'Sunday')
AND shipped_date IS NOT NULL
ORDER BY order_date DESC;
```

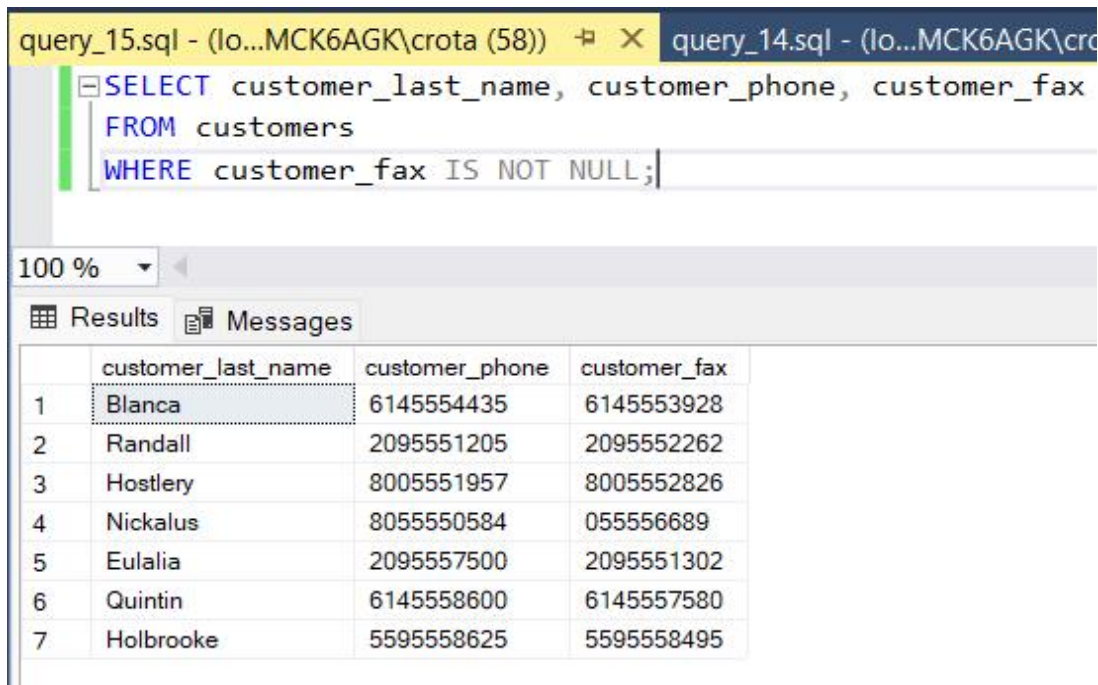
Below the query editor, the results are displayed in a table with 5 columns: order_id, order_date, shipped_date, and Order_Day. The results are sorted by order_date in descending order.

	order_id	order_date	shipped_date	Order_Day
1	796	2014-07-19	2014-07-26	Saturday
2	778	2014-07-12	2014-07-21	Saturday
3	693	2014-06-07	2014-06-19	Saturday
4	624	2014-05-04	2014-05-09	Sunday
5	550	2014-03-23	2014-04-03	Sunday
6	548	2014-03-22	2014-04-18	Saturday
7	491	2014-02-08	2014-02-14	Saturday
8	442	2013-12-28	2014-01-03	Saturday
9	298	2013-08-18	2013-09-22	Sunday
10	118	2013-02-24	2013-02-28	Sunday
11	89	2013-01-20	2013-01-22	Sunday
12	45	2012-11-25	2012-11-30	Sunday
13	32	2012-11-10	2012-11-13	Saturday

Explanation: Selecting the identifying column for orders, the date of orders, the date an order was shipped, and the day of the week the order was placed as the fields of interest, all derived or calculated from the table of orders. To get the name of the day, the date name function is used which takes a time unit and then the desired date as inputs. Only orders that have been shipped will be returned in the query results and sorted by most recent to later order dates.

Query 15: Write a query to display the customer's last name, phone number, and fax number but only display those customers that have a fax number.

```
SELECT customer_last_name, customer_phone, customer_fax
FROM customers
WHERE customer_fax IS NOT NULL;
```



The screenshot shows a SQL query editor with the following query:

```
SELECT customer_last_name, customer_phone, customer_fax
FROM customers
WHERE customer_fax IS NOT NULL;
```

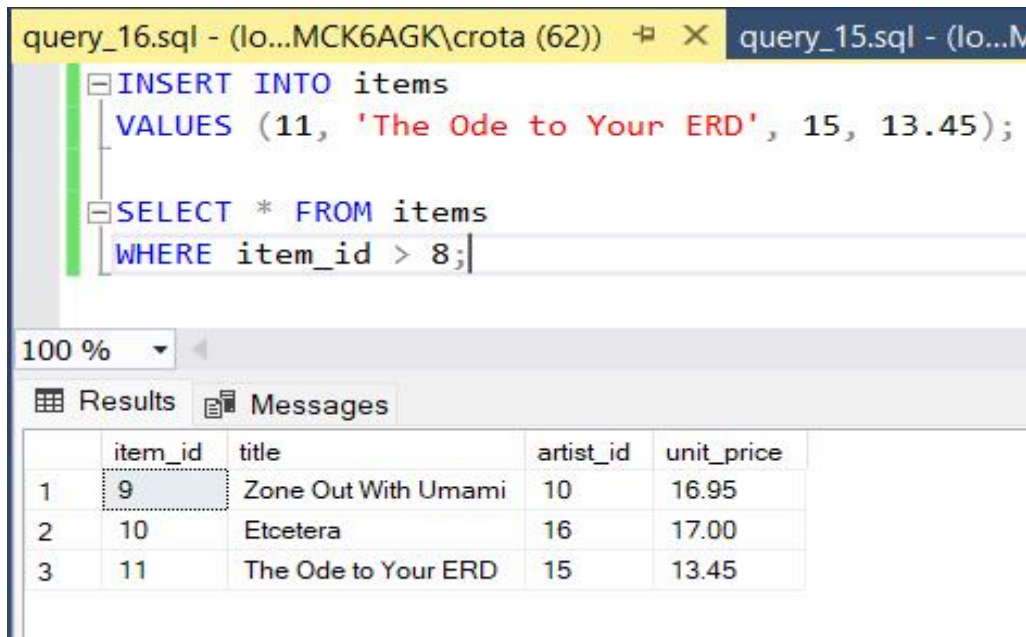
Below the query editor, the results window is displayed, showing a table with 7 rows of data. The columns are customer_last_name, customer_phone, and customer_fax.

	customer_last_name	customer_phone	customer_fax
1	Blanca	6145554435	6145553928
2	Randall	2095551205	2095552262
3	Hostlery	8005551957	8005552826
4	Nickalus	8055550584	055556689
5	Eulalia	2095557500	2095551302
6	Quintin	6145558600	6145557580
7	Holbrooke	5595558625	5595558495

Explanation: The query displays the three column fields with the customer's last name, their phone number, and their fax number which is derived from the customers table. Because we want to see all customers who actually have a fax number, customers who do not have a fax number should not be shown/returned. By putting the condition where the customer fax field must not be empty (null) then we achieve this effect.

Query 16: Create a statement to insert a new record into the items table with the following values: item_id: 11, title: The Ode to Your ERD, Artist_id: 15, unit_price: 13.45

```
INSERT INTO items VALUES (11, 'The Ode to Your ERD', 15, 13.45); SELECT *  
FROM items WHERE item_id > 8;
```



```
query_16.sql - (lo...MCK6AGK\crota (62)) X query_15.sql - (lo...M
```

```
INSERT INTO items  
VALUES (11, 'The Ode to Your ERD', 15, 13.45);  
  
SELECT * FROM items  
WHERE item_id > 8;
```

100 %

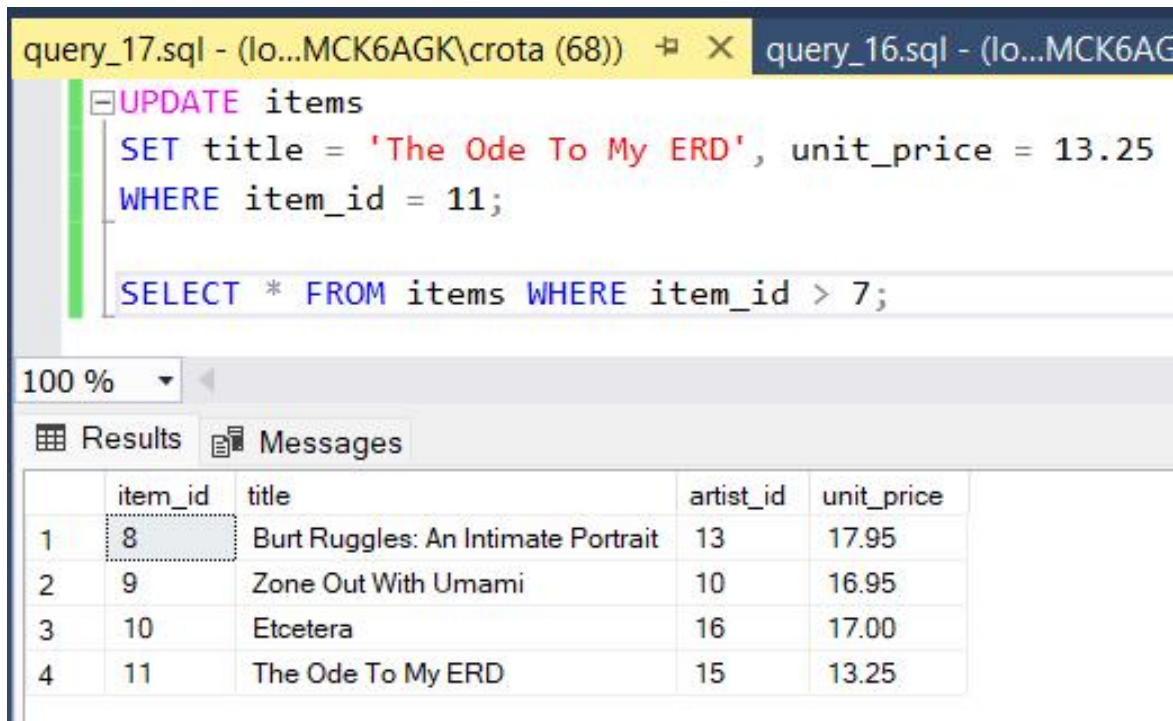
Results Messages

	item_id	title	artist_id	unit_price
1	9	Zone Out With Umami	10	16.95
2	10	Etcetera	16	17.00
3	11	The Ode to Your ERD	15	13.45

Explanation: Using the insert function we specify the table that the values should be inserted into which in this case is the items table. These new values are placed by order of the columns of the items table as seen in the parenthesis as the last row in the items table. The select statement is then used to verify that a new row had been created after item 10.

Query 17: Create a statement to update the record inserted in the previous step to change the unit price of this item to 13.25 with item_id: 11, title: The Ode To My ERD, artist: 15, and unit_price: 13.25.

```
UPDATE items SET title = 'The Ode To My ERD', unit_price = 13.25  
WHERE item_id = 11; SELECT * FROM items WHERE item_id > 7;
```



The screenshot shows a SQL query editor with two tabs: 'query_17.sql' and 'query_16.sql'. The active tab 'query_17.sql' contains the following SQL code:

```
UPDATE items  
SET title = 'The Ode To My ERD', unit_price = 13.25  
WHERE item_id = 11;  
  
SELECT * FROM items WHERE item_id > 7;
```

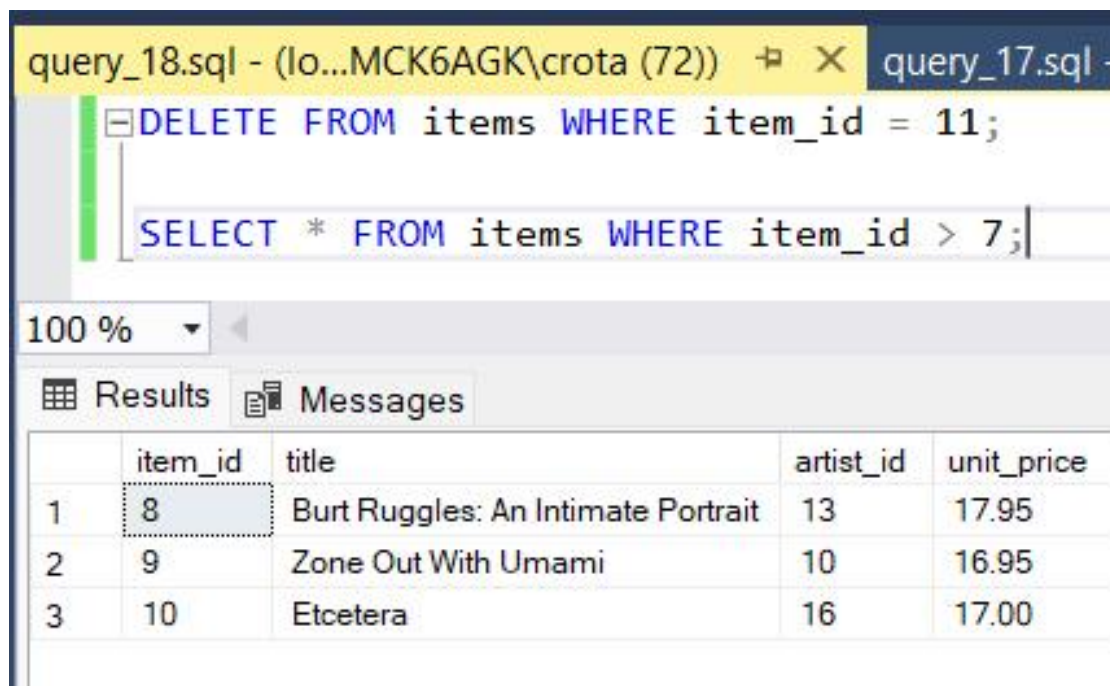
Below the query editor, there is a 'Results' tab showing the output of the query. The results are displayed in a table with 5 columns: 'item_id', 'title', 'artist_id', and 'unit_price'. The table contains 4 rows of data, with the first row (item_id 8) highlighted.

	item_id	title	artist_id	unit_price
1	8	Burt Ruggles: An Intimate Portrait	13	17.95
2	9	Zone Out With Umami	10	16.95
3	10	Etcetera	16	17.00
4	11	The Ode To My ERD	15	13.25

Explanation: In query 16 the inserted row had “The Ode to Your ERD” as the title and “13.25” as the unit price. In this query the update clause is used to update the table items with the condition that the update will take place at item 11 and the new values of “The Ode to My ERD” is set to the title and the unit price of 13.25. The select clause is used to verify that the row had been successfully updated.

Query 18: Create a statement to delete the entire record that was inserted and then updated in the previous steps. Show your DELETE statement along with the results of the following SELECT query to verify that the insert worked correctly.

```
DELETE FROM items WHERE item_id = 11;  
SELECT * FROM items WHERE item_id > 7;
```



The screenshot shows a SQL query editor with two tabs: 'query_18.sql' and 'query_17.sql'. The 'query_18.sql' tab is active and contains the following SQL code:

```
DELETE FROM items WHERE item_id = 11;  
  
SELECT * FROM items WHERE item_id > 7;
```

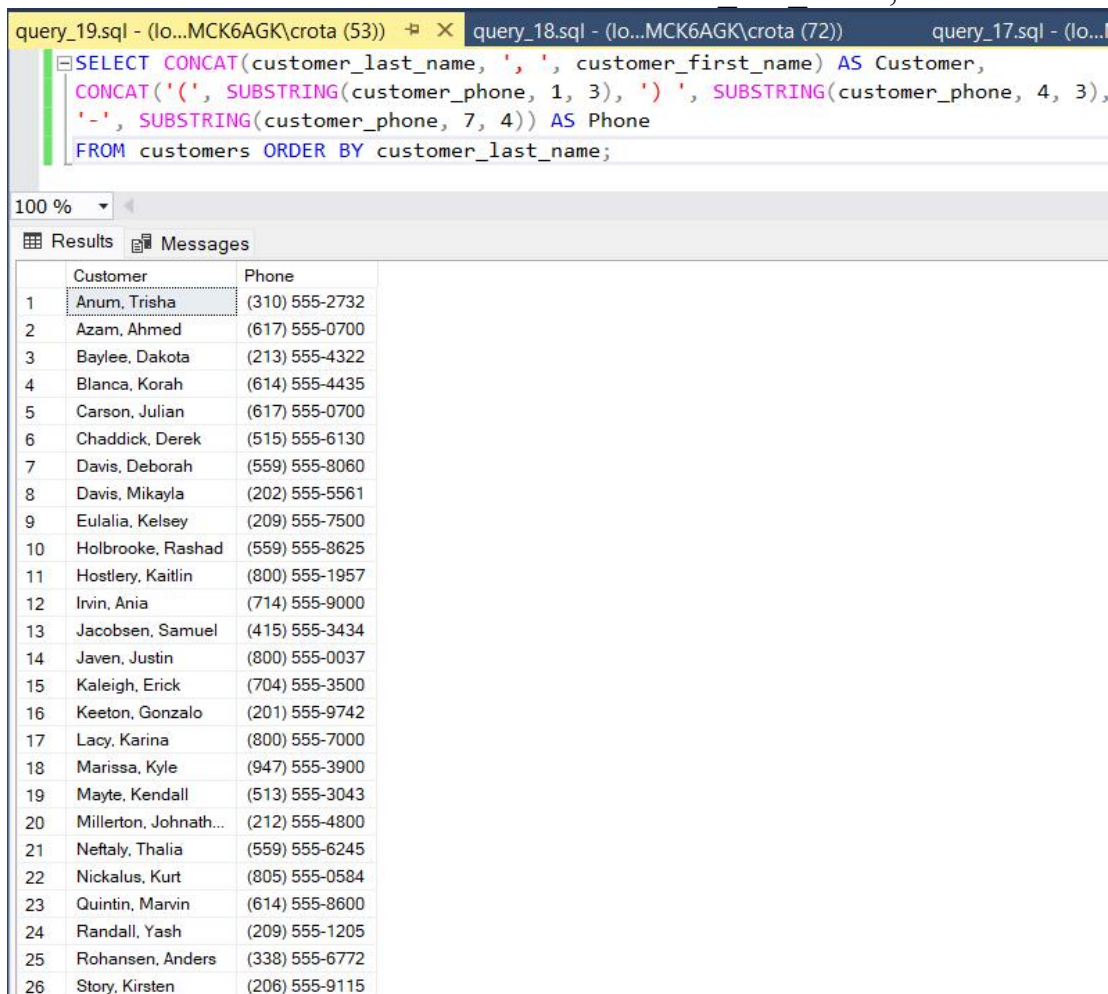
Below the code editor, there is a 'Results' tab and a 'Messages' tab. The 'Results' tab is selected and displays the results of the SELECT query. The results are shown in a table with the following columns: 'item_id', 'title', 'artist_id', and 'unit_price'.

	item_id	title	artist_id	unit_price
1	8	Burt Ruggles: An Intimate Portrait	13	17.95
2	9	Zone Out With Umami	10	16.95
3	10	Etcetera	16	17.00

Explanation: Using the delete clause, the new row inserted in query 16 is now deleted when specified the table items from which the row is in and the condition that the row has the value 11 for its item id to signify we are removing item 11. Then the select statement is used to show all rows after item 7 to see that the row with item 11 has been removed.

Query 19: Using the SUBSTRING and CONCAT functions, write a query to display each customer name as a single field in the format “Jones, Tom” with a heading of Customer along with the customer_phone field in a nicely formatted calculated column named Phone. For example, a record containing the customer_phone value 9095595443 would be output with parentheses, spaces, and hyphens, like this: (909) 559-5443. Sort by last name.

```
SELECT CONCAT(customer_last_name, ', ', customer_first_name) AS Customer,
CONCAT('(', SUBSTRING(customer_phone, 1, 3), ') ',
SUBSTRING(customer_phone, 4, 3), '-', SUBSTRING(customer_phone, 7, 4)) AS Phone
FROM customers ORDER BY customer_last_name;
```



The screenshot shows a SQL query editor with the following query:

```
SELECT CONCAT(customer_last_name, ', ', customer_first_name) AS Customer,
CONCAT('(', SUBSTRING(customer_phone, 1, 3), ') ',
SUBSTRING(customer_phone, 4, 3), '-', SUBSTRING(customer_phone, 7, 4)) AS Phone
FROM customers ORDER BY customer_last_name;
```

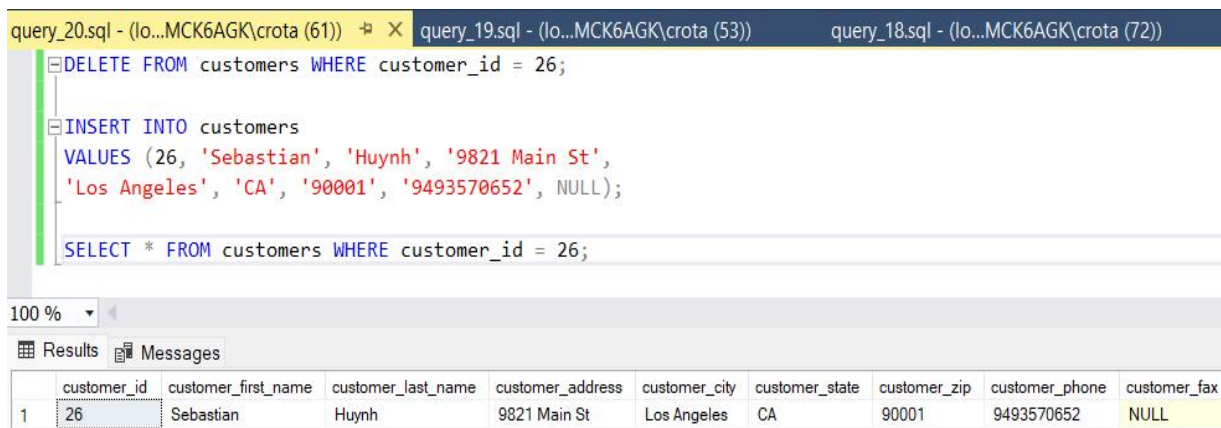
The results are displayed in a table with two columns: Customer and Phone. The results are sorted by last name.

	Customer	Phone
1	Anum, Trisha	(310) 555-2732
2	Azam, Ahmed	(617) 555-0700
3	Baylee, Dakota	(213) 555-4322
4	Blanca, Korah	(614) 555-4435
5	Carson, Julian	(617) 555-0700
6	Chaddick, Derek	(515) 555-6130
7	Davis, Deborah	(559) 555-8060
8	Davis, Mikayla	(202) 555-5561
9	Eulalia, Kelsey	(209) 555-7500
10	Holbrooke, Rashad	(559) 555-8625
11	Hostlery, Kaitlin	(800) 555-1957
12	Ivin, Ania	(714) 555-9000
13	Jacobsen, Samuel	(415) 555-3434
14	Javen, Justin	(800) 555-0037
15	Kaleigh, Erick	(704) 555-3500
16	Keeton, Gonzalo	(201) 555-9742
17	Lacy, Karina	(800) 555-7000
18	Marissa, Kyle	(947) 555-3900
19	Mayte, Kendall	(513) 555-3043
20	Millerton, Johnath...	(212) 555-4800
21	Neftaly, Thalia	(559) 555-6245
22	Nickalus, Kurt	(805) 555-0584
23	Quintin, Marvin	(614) 555-8600
24	Randall, Yash	(209) 555-1205
25	Rohansen, Anders	(338) 555-6772
26	Story, Kirsten	(206) 555-9115

Explanation: Derived from customers table. Combining two columns with the format of “last name, first name” for each customer and the alias given to this field is “Customer”. The other column is a formatted string converting raw phone numbers into the format seen with parenthesis and semicolon. This is done with the concatenating different portions of the phone number using substring and adding the appropriate characters together. The results are then sorted by last name.

Query 20: Create a statement to insert a new record with your values: your customer id, first name, last name, address, city, state, zip code and fax number. Use your real name and (9821 Main St, Los Angeles, CA 90001) as your address: Show your INSERT statement along with the results of the following SELECT query to verify that the insert worked correctly.

```
DELETE FROM customers WHERE customer_id = 26;  
INSERT INTO customers VALUES (26, 'Sebastian', 'Huynh', '9821 Main St', 'Los Angeles', 'CA', '90001', '9493570652', NULL);  
SELECT * FROM customers WHERE customer_id = 26;
```



```
query_20.sql - (lo...MCK6AGK\crota (61)) X query_19.sql - (lo...MCK6AGK\crota (53)) query_18.sql - (lo...MCK6AGK\crota (72))  
DELETE FROM customers WHERE customer_id = 26;  
  
INSERT INTO customers  
VALUES (26, 'Sebastian', 'Huynh', '9821 Main St',  
      'Los Angeles', 'CA', '90001', '9493570652', NULL);  
  
SELECT * FROM customers WHERE customer_id = 26;
```

100 %

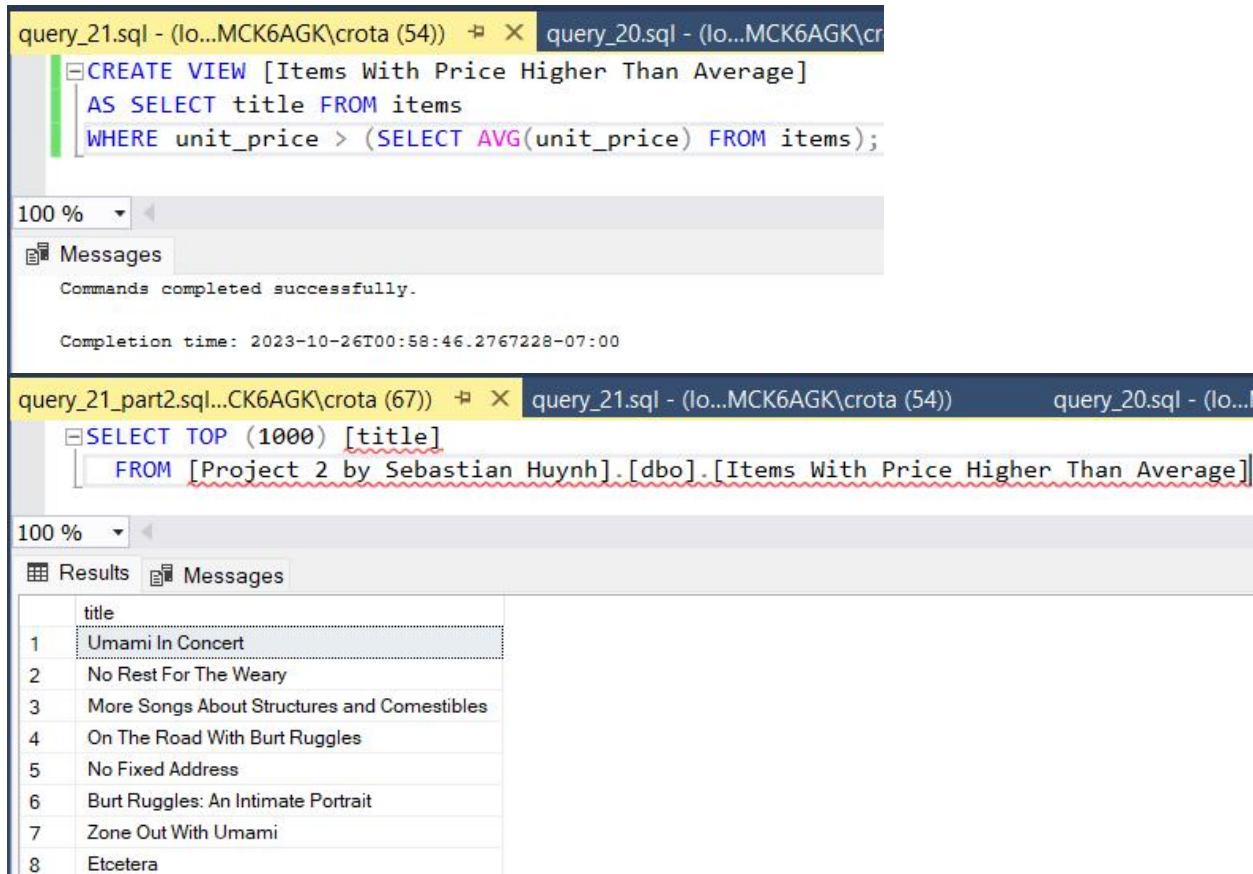
Results Messages

	customer_id	customer_first_name	customer_last_name	customer_address	customer_city	customer_state	customer_zip	customer_phone	customer_fax
1	26	Sebastian	Huynh	9821 Main St	Los Angeles	CA	90001	9493570652	NULL

Explanation: Adding my own information as if I was a customer. Customer 26 already exists so in order to add my information without throwing an error, the delete clause for the table of customers is used. Then the insert clause is used with my values placed in parenthesis in the order of the columns of the customer table. The select clause is used to display that the row has been updated.

Query 21: Creates a view that selects every title in the "item" table with a price higher than the average price.

```
CREATE VIEW [Items With Price Higher Than Average] AS SELECT title FROM items WHERE unit_price > (SELECT AVG(unit_price) FROM items);  
SELECT TOP (1000) [title] FROM [Project 2 by Sebastian Huynh].[dbo].[Items With Price Higher Than Average]
```

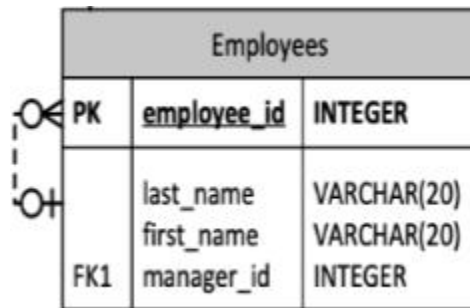


The screenshot shows two queries being executed in SQL Server Enterprise Manager. The first query, 'query_21.sql', creates a view named 'Items With Price Higher Than Average' by selecting titles from the 'items' table where the unit price is greater than the average unit price. The second query, 'query_21_part2.sql', selects the top 1000 titles from this view. The results of the second query are displayed in a table below.

	title
1	Umami In Concert
2	No Rest For The Weary
3	More Songs About Structures and Comestibles
4	On The Road With Burt Ruggles
5	No Fixed Address
6	Burt Ruggles: An Intimate Portrait
7	Zone Out With Umami
8	Etcetera

Explanation: Created a view (specialized table object) that contains one column field from the items table which are the item titles. “Items with a price higher than average” is the name for this view. The avg function is used on the unit prices to be the condition for the query and unit prices must be larger than it to be returned. Finally, the last select clause was autogenerated by management studio when right clicking ‘select top 1000 rows’ for the view in the SSMS object explorer. This is a quick way to see the first 1000 rows of a table or view to double check work.

Query 22: Explain the cardinality from the employees-to-employees table.



- Explanation: The cardinality of the employees table is a non-identifying unary relationship. It is non-identifying because the primary key is not a composite key made up of the foreign key and primary key. It is unary because managers are also employees with a different responsibility. According to the relationship depicted above, a manager may have zero or many employees under him/her but an employee may only belong to zero or one manager. A manager's primary key will be his/her employee id but it will simultaneously appear as the foreign key (manager id) for the employee(s) under him/her.

Query 23: Explain the following statement.

```
SELECT pv.ProductID, v.BusinessEntityID, v.Name
FROM Purchasing.ProductVendor AS pv
INNER JOIN Purchasing.Vendor AS v
ON (pv.BusinessEntityID = v.BusinessEntityID)
WHERE StandardPrice > $10
AND Name LIKE N'F%';
```

- Explanation: This query statement inner/equi joins two tables. We first want this combined table to have the column fields of product id from the ProductVendor table and then the business entity id and name from the Vendor table. Both tables are part of the Purchasing schema of the database and are given alias names with ProductVendor given the name pv and Vendor given the name v. The tables are joined on the common column of business entity id and will only return matching rows that fit the description of having a standard price greater than ten dollars and any length name that starts with letter F.

Query 24: Explain the following statement.

```
SELECT pv.ProductID, v.BusinessEntityID, v.Name
FROM Purchasing.ProductVendor AS pv, Purchasing.Vendor AS v
WHERE pv.BusinessEntityID=v.BusinessEntityID
      AND StandardPrice > $10
      AND Name LIKE N'F%';
```

- Explanation: This query statement inner/equijoins two tables by use of the where clause. This combined table has the column fields of product id from the ProductVendor table and then the business entity id and name from the Vendor table. Both tables are part of the Purchasing schema of the database and are given alias names with ProductVendor given the name pv and Vendor given the name v. The tables are joined on the common column of business entity id and will only return matching rows that fit the description of having a standard price greater than ten dollars and any length name that starts with the letter F.

Query 25: Compare query 23 and query 24.

- Explanation: Both queries output the same result but the order of execution is different for them. Join clauses usually execute before any of the where clauses so the join would combine the two tables first in query 23 then filter them with where. But in query 24 the join occurs when the where clause filters out all the rows derived using the from clause to match conditions.

Results & Discussion

By practicing along with the project guidelines, SQL is a powerful language that makes relational databases useful to businesses. This is seen with how helpful queries can be, in the sense that they can show us what a business wants to see quickly rather than sifting through numerous rows and columns manually. This “narrowing down” effect for finding specific data lessens the chance for human error to be made and also highlights any useful connections/relationships in data that may not be easily seen with everything presented together. A drawback to writing SQL is how rigid the rules can be, meaning things that may make sense to you will not make sense to the program or database if syntax is missing or placed incorrectly or disobeys a SQL rule. This happened to me on various occasions during the completion of the project. For example, when trying to reference an expression using its alias name in the where clause throws an error, so I had to repeat the entire expression again for the query to run. Another example is when I attempted to use the like clause and accidentally put “is like” which makes sense from the English way of structuring a sentence, but SQL doesn’t accept it.

Lessons Learned

1. Just because the SQL code runs does not mean that the results returned are what I intended it to show.
 - In query #1, initially I did not put ASC/ascending for the second column in the “order by” clause and it did not sort the second column even though it sorted by the first column. So it is important to always check if data corresponds to intention.
2. Formatting is important, especially if a company establishes how data should be presented or formatting makes data realistic.
 - When calculating the discounted price in query #7, having many decimal places in the answer is not useful because cents only go up to two decimal places. Also extra zeroes at the end are not helpful and can confuse those looking at the data which is why the convert to double precision formula was used.
3. There are many possible ways to get the same result.
 - Like observed in query 25
 - Concatenating strings can be done using “concat” or by simply using the “+” sign.

Conclusion

Since the 1970s, SQL has proven to be one of the best ways to manage relational databases where collaboration between users, programs, and databases is made possible. However, there is also NoSql which stands for “not only sql” and it has been gaining a lot more popularity in recent years. With the rise of many unconventional types of data and internet-based applications, “NoSql databases thrive in environments with highly dynamic data and scale much better than SQL databases” (Powell, 2021).

Cited References

1. Hoffer, Jeff, et al. Modern Database Management. Available from: VitalSource Bookshelf, (13th Edition). Pearson Education (US), 2018.
2. Panigrahi, Kiran Kumar. “Difference between DDL and DML in DBMS.” Online Tutorials, Courses, and eBooks Library, 12 Sept. 2023, www.tutorialspoint.com/difference-between-ddl-and-dml-in-dbms.
3. Powell, Ron. “SQL VS NoSQL Databases.” CircleCI, CircleCI, 10 Sept. 2021, circleci.com/blog/sql-vs-nosql-databases/.